

SQL Quick Reference Guide

For Kernel Module Detective CTF

■ Table of Contents

1. Basic SELECT Queries
2. Filtering with WHERE
3. Sorting and Limiting
4. Aggregate Functions
5. GROUP BY and HAVING
6. JOIN Operations
7. Subqueries
8. Common Table Expressions (CTEs)
9. Set Operations (UNION, INTERSECT, EXCEPT)
10. Window Functions
11. Operator Precedence
12. SQLite-Specific Features

1. Basic SELECT Queries

Retrieve data from a single table:

Select all columns:

```
SELECT * FROM table_name;
```

Select specific columns:

```
SELECT column1, column2, column3  
FROM table_name;
```

Select with column aliases:

```
SELECT column1 AS alias1, column2 AS alias2  
FROM table_name;
```

Select distinct values:

```
SELECT DISTINCT column_name  
FROM table_name;
```

2. Filtering with WHERE

Filter rows based on conditions:

Basic comparison operators:

```
SELECT * FROM table_name  
WHERE column1 = 'value';  
  
-- Operators: =, !=, <, >, <=, >=, <> (not equal)
```

Multiple conditions (AND, OR):

```
SELECT * FROM table_name  
WHERE condition1 AND condition2;  
  
SELECT * FROM table_name  
WHERE condition1 OR condition2;
```

IN operator:

```
SELECT * FROM table_name  
WHERE column_name IN ('value1', 'value2', 'value3');
```

BETWEEN operator:

```
SELECT * FROM table_name  
WHERE column_name BETWEEN 100 AND 200;
```

LIKE operator (pattern matching):

```
SELECT * FROM table_name  
WHERE column_name LIKE 'pattern%';  
  
-- % matches any sequence  
-- _ matches single character
```

NULL checks:

```
SELECT * FROM table_name  
WHERE column_name IS NULL;  
  
SELECT * FROM table_name  
WHERE column_name IS NOT NULL;
```

3. Sorting and Limiting Results

ORDER BY (ascending):

```
SELECT * FROM table_name  
ORDER BY column1 ASC;
```

ORDER BY (descending):

```
SELECT * FROM table_name  
ORDER BY column1 DESC;
```

Multiple sort columns:

```
SELECT * FROM table_name  
ORDER BY column1 DESC, column2 ASC;
```

LIMIT results:

```
SELECT * FROM table_name  
LIMIT 10;
```

LIMIT with OFFSET:

```
SELECT * FROM table_name  
LIMIT 10 OFFSET 20; -- Skip first 20, return next 10
```

4. Aggregate Functions

Perform calculations on sets of rows:

Common aggregate functions:

```
SELECT COUNT(*) FROM table_name;
SELECT COUNT(DISTINCT column) FROM table_name;

SELECT SUM(column_name) FROM table_name;
SELECT AVG(column_name) FROM table_name;
SELECT MIN(column_name) FROM table_name;
SELECT MAX(column_name) FROM table_name;
```

Combining aggregates:

```
SELECT
COUNT(*) AS total_rows,
AVG(column1) AS average_value,
MAX(column2) AS max_value
FROM table_name;
```

5. GROUP BY and HAVING

Basic GROUP BY:

```
SELECT column1, COUNT(*)
FROM table_name
GROUP BY column1;
```

Multiple grouping columns:

```
SELECT column1, column2, COUNT(*)
FROM table_name
GROUP BY column1, column2;
```

HAVING clause (filter after grouping):

```
SELECT column1, COUNT(*) AS count
FROM table_name
GROUP BY column1
HAVING COUNT(*) > 5;
```

HAVING vs WHERE:

- WHERE filters rows BEFORE grouping
- HAVING filters groups AFTER aggregation

Example combining WHERE and HAVING:

```
SELECT module_name, COUNT(*) AS failures
FROM module_events
WHERE status = 'FAILED'
GROUP BY module_name
HAVING COUNT(*) > 3
ORDER BY failures DESC;
```

6. JOIN Operations

Combine data from multiple tables:

INNER JOIN (matching rows only):

```
SELECT t1.column1, t2.column2
FROM table1 AS t1
INNER JOIN table2 AS t2
ON t1.id = t2.foreign_id;
```

LEFT JOIN (all rows from left table):

```
SELECT t1.column1, t2.column2
FROM table1 AS t1
LEFT JOIN table2 AS t2
ON t1.id = t2.foreign_id;
```

Multiple JOINS:

```
SELECT t1.col1, t2.col2, t3.col3
FROM table1 AS t1
INNER JOIN table2 AS t2 ON t1.id = t2.id
INNER JOIN table3 AS t3 ON t2.id = t3.id
WHERE t1.status = 'active';
```

Self JOIN:

```
SELECT a.column1, b.column2
FROM table AS a
INNER JOIN table AS b
ON a.id = b.parent_id;
```

JOIN with aggregation:

```
SELECT t1.module_name, COUNT(t2.error_id) AS error_count
FROM module_events AS t1
LEFT JOIN error_codes AS t2
ON t1.module_name = t2.affected_module
GROUP BY t1.module_name;
```

7. Subqueries

Subquery in WHERE clause:

```
SELECT * FROM table1
WHERE column1 IN (
  SELECT column2 FROM table2
  WHERE condition
);
```

Subquery in FROM clause:

```
SELECT subq.col1, subq.count
FROM (
  SELECT column1 AS col1, COUNT(*) AS count
  FROM table_name
  GROUP BY column1
) AS subq
WHERE subq.count > 10;
```

Correlated subquery:

```
SELECT *
FROM table1 AS t1
WHERE EXISTS (
  SELECT 1 FROM table2 AS t2
  WHERE t2.foreign_id = t1.id
  AND t2.status = 'FAILED'
);
```

Subquery with aggregates:

```
SELECT *
FROM table1
WHERE value > (
  SELECT AVG(value) FROM table1
);
```

8. Common Table Expressions (CTEs)

Named temporary result sets for complex queries:

Basic CTE:

```
WITH cte_name AS (
  SELECT column1, COUNT(*) AS count
  FROM table_name
  GROUP BY column1
)
SELECT * FROM cte_name
WHERE count > 5;
```

Multiple CTEs:

```
WITH
failed_modules AS (
  SELECT module_name, COUNT(*) AS failures
  FROM module_events
  WHERE status = 'FAILED'
  GROUP BY module_name
),
critical_errors AS (
  SELECT affected_module, COUNT(*) AS errors
  FROM error_codes
  WHERE severity = 'CRITICAL'
  GROUP BY affected_module
)
SELECT
fm.module_name,
fm.failures,
ce.errors
FROM failed_modules AS fm
INNER JOIN critical_errors AS ce
ON fm.module_name = ce.affected_module;
```

Recursive CTE (advanced):

```
WITH RECURSIVE counter(n) AS (
  SELECT 1
  UNION ALL
  SELECT n + 1 FROM counter WHERE n < 10
)
SELECT * FROM counter;
```

9. Set Operations

UNION (combine, remove duplicates):

```
SELECT column1 FROM table1
UNION
SELECT column1 FROM table2;
```

UNION ALL (combine, keep duplicates):

```
SELECT column1 FROM table1
UNION ALL
SELECT column1 FROM table2;
```

INTERSECT (common rows):

```
SELECT module_name FROM module_events
WHERE status = 'FAILED'
INTERSECT
SELECT affected_module FROM error_codes
WHERE severity = 'CRITICAL';
```

EXCEPT (rows in first but not second):

```
SELECT module_name FROM module_events
EXCEPT
SELECT module_name FROM module_events
WHERE status = 'SUCCESS';
```

■■ Note: Column count and types must match in all queries

10. Window Functions

Perform calculations across rows without collapsing results:

ROW_NUMBER:

```
SELECT
module_name,
failures,
ROW_NUMBER() OVER (ORDER BY failures DESC) AS rank
FROM module_summary;
```

RANK and DENSE_RANK:

```
SELECT
module_name,
error_count,
RANK() OVER (ORDER BY error_count DESC) AS rank,
DENSE_RANK() OVER (ORDER BY error_count DESC) AS dense_rank
FROM error_summary;
```

Partition by (grouping within windows):

```
SELECT
boot_session,
module_name,
COUNT(*) AS errors,
ROW_NUMBER() OVER (
PARTITION BY boot_session
ORDER BY COUNT(*) DESC
) AS rank_in_session
FROM error_codes
GROUP BY boot_session, module_name;
```

Aggregate window functions:

```
SELECT
timestamp,
value,
AVG(value) OVER (
ORDER BY timestamp
ROWS BETWEEN 2 PRECEDING AND CURRENT ROW
) AS moving_avg
FROM measurements;
```

11. Operator Precedence

Understanding how SQL evaluates complex conditions:

Precedence order (highest to lowest):

1	Parentheses ()	Force evaluation order
2	Comparison	=, !=, <, >, <=, >=, IN, BETWEEN, LIKE
3	NOT	Logical NOT
4	AND	Logical AND
5	OR	Logical OR

Example WITHOUT parentheses:

```
-- Evaluates as: status = 'FAILED' AND (severity = 'HIGH' OR severity = 'CRITICAL')
SELECT * FROM errors
WHERE status = 'FAILED' AND severity = 'HIGH' OR severity = 'CRITICAL';
```

Example WITH parentheses (recommended):

```
-- Explicit grouping makes intent clear
SELECT * FROM errors
WHERE status = 'FAILED'
AND (severity = 'HIGH' OR severity = 'CRITICAL');
```

Complex example:

```
SELECT * FROM events
WHERE (type = 'ERROR' OR type = 'WARNING')
AND module_name IN ('mod1', 'mod2')
AND (timestamp > 1000 OR priority = 'HIGH')
AND NOT ignored = 1;
```

12. SQLite-Specific Features

Database metadata:

```
-- List all tables
SELECT name FROM sqlite_master
WHERE type='table';

-- Get table schema
PRAGMA table_info(table_name);
```

Type conversion (CAST):

```
SELECT CAST(column AS INTEGER) FROM table;
SELECT CAST(column AS REAL) FROM table;
SELECT CAST(column AS TEXT) FROM table;
```

String functions:

```
SELECT LENGTH(column) FROM table;
SELECT UPPER(column) FROM table;
SELECT LOWER(column) FROM table;
SELECT SUBSTR(column, 1, 5) FROM table;
SELECT REPLACE(column, 'old', 'new') FROM table;
```

Math functions:

```
SELECT ABS(column) FROM table;
SELECT ROUND(column, 2) FROM table;
SELECT MAX(col1, col2) FROM table;
SELECT MIN(col1, col2) FROM table;
```

Date/time functions:

```
SELECT datetime(timestamp, 'unixepoch') FROM table;
SELECT strftime('%Y-%m-%d', timestamp, 'unixepoch') FROM table;
```

EXPLAIN QUERY PLAN (analyze performance):

```
EXPLAIN QUERY PLAN
SELECT * FROM table WHERE column = 'value';
```

■ Quick Tips & Best Practices

Use aliases: Make complex queries readable with table and column aliases

Indent properly: Format multi-line queries with proper indentation

Test incrementally: Build complex queries step by step

Use EXPLAIN: Check query execution plans for performance

Comment your code: Use -- for single line or /* */ for blocks

Be explicit: Use parentheses in complex WHERE clauses

Choose right JOIN: INNER for matches, LEFT for optional relationships

GROUP before FILTER: WHERE filters rows, HAVING filters groups

Name your CTEs: Use descriptive names for Common Table Expressions

Watch for NULL: Use IS NULL, not = NULL

Good luck with the CTF! ■